




Five agent designs and their unchecked authority

Each class gets work done. Each keeps some of five authorities (close, record, share, remember, improve) in the model that does the work.

C1. Direct loop	C2. Static workflow graph	C3. Role ensemble	C4. Search and evolution	C5. Memory-augmented
one model in a think-act-observe loop; it submits when it says it is done	human-authored DAG of LLM steps; edges fixed before the run begins	copies of one model debate, vote, or play roles	tree search or prompt evolution scored by the model itself	retrieval store or summaries re-injected into the prompt
				
kept in-model	kept in-model	kept in-model	kept in-model	kept in-model
close record share remember improve	close record share, remember, improve: no mechanism at all	close share record remember, improve: no mechanism at all	close improve record, share, remember: no mechanism at all	remember close record, share, improve: no mechanism at all
where it breaks The model grades its own work. Fluent failures cross as passes.	where it breaks Off-script states have no route. The judging node is the same model kind.	where it breaks Correlated errors vote together. Consensus is not correctness.	where it breaks Search optimizes the model's own confidence. Cost grows per branch.	where it breaks Unvalidated writes persist. Retrieval relevance is not correctness.

Shared root cause

In every class the component that produces the work also holds authority over it: it declares the result, owns the history, writes shared state, decides what persists, or promotes its own changes. Model confidence substitutes for execution truth, and every added safeguard is one more opinion from the same source.

Five authorities move out of the model and into a replayable substrate.

Class evidence is cited in the bibliography; the argument here is about shapes, not products.