



Five ways agent harnesses are built today, and the authority each leaves inside the model

Each class gets work done. Each keeps some of five authorities (close, record, share, remember, improve) in the model that does the work.

C1. Direct loop	C2. Static workflow graph	C3. Role ensemble	C4. Search and evolution	C5. Memory-augmented
one model in a think-act-observe loop; it submits when it says it is done  kept in-model close record share remember improve where it breaks The model grades its own work. Fluent failures cross as passes.	human-authored DAG of LLM steps; edges fixed before the run begins  kept in-model close record <i>share, remember, improve: no mechanism at all</i> where it breaks Off-script states have no route. The judging node is the same model kind.	copies of one model debate, vote, or play roles  kept in-model close share record <i>remember, improve: no mechanism at all</i> where it breaks Correlated errors vote together. Consensus is not correctness.	tree search or prompt evolution scored by the model itself  kept in-model close improve <i>record, share, remember: no mechanism at all</i> where it breaks Search optimizes the model's own confidence. Cost grows per branch.	retrieval store or summaries re-injected into the prompt  kept in-model remember close <i>record, share, improve: no mechanism at all</i> where it breaks Unvalidated writes persist. Retrieval relevance is not correctness.

Shared root cause

In every class the component that produces the work also holds authority over it: it declares the result, owns the history, writes shared state, decides what persists, or promotes its own changes. Model confidence substitutes for execution truth, and every added safeguard is one more opinion from the same source.

This paper moves the five authorities out of the model and into a replayable substrate. Figure 2 shows the resulting pipeline.
Class evidence is cited in the bibliography; the argument here is about shapes, not products.